



Time-correct memory for *AI* agents

Open-source infrastructure for point-in-time recall, provenance, and auditable agent decisions.

AI agents answer historical questions with facts they learned later



Wrong-time recall

A current value silently replaces the value that was true when the decision was made.



Missing evidence

Teams cannot reliably show which source, version, or policy supported an answer.



Costly reconstruction

Audits become manual investigations across changing documents, tables, and memories.

Lians makes agent memory reconstructable by design

Bitemporal facts

Track when a fact was true and when the system learned it.

Provenance

Preserve the source and history behind every remembered fact.


Point-in-time recall

Ask what the agent could have known at any prior moment.

Deterministic change

Handle supersession, corrections, and erasure with explicit rules.

One memory layer turns changing facts into defensible answers



The first wedge is regulated, high-consequence agent workflows

INITIAL USERS

Teams building AI agents for financial research, risk, compliance, governance, and enterprise operations.
They need historical truth and evidence, not only semantic retrieval.

Regulated enterprises

Audit-ready recall
Policy and source history
Private deployment paths

AI infrastructure teams

Memory SDK integration
Versioned retrieval
Open-source foundation

The product is live, open, and ready for design partners

CURRENT EVIDENCE

We launched the initial product in July 2026. We are pre-revenue and have not raised outside capital.

Apache 2.0

Open-source
repository

Public demo

Product can be
evaluated today

Pre-revenue

No invented customer
claims

Open source earns adoption; enterprise delivery creates revenue

Free core

Apache 2.0 SDK and core temporal-memory capabilities.

Enterprise

Private deployment, security hardening, support, and SLAs.

Managed platform

Hosted operation, observability, and production reliability.

Services

Integrations, implementation, and evidence workflows.



Help AI agents remember what was true then

lians.ai | github.com/Lians-ai/Lians

info@lians.ai

Ethan Beirne + Dereck
Salazar